IBM watsonx.data

Seller guidance and legal disclaimer

IBM and Business Partner
Internal Use Only

Slides in this presentation marked as **"IBM and Business Partner Internal Use Only"** are for IBM and Business Partner use and should not be shared with clients or anyone else outside of IBM or the Business Partners' company.

The information contained in this publication is provided for informational purposes only. While efforts were made to verify the completeness and accuracy of the information contained in this publication, it is provided AS IS without warranty of any kind, express or implied. In addition, this information is based on IBM's current product plans and strategy, which are subject to change by IBM without notice. IBM shall not be responsible for any damages arising out of the use of, or otherwise related to, this publication or any other materials. Nothing contained in this publication is intended to, nor shall have the effect of, creating any warranties or representations from IBM or its suppliers or licensors, or altering the terms and conditions of the applicable license agreement governing the use of IBM software.

References in this presentation to IBM products, programs, or services do not imply that they will be available in all countries in which IBM operates. Product release dates and/or capabilities referenced in this presentation may change at any time at IBM's sole discretion based on market opportunities or other factors and are not intended to be a commitment to future product or feature availability in any way. Nothing contained in these materials is intended to, nor shall have the effect of, stating or implying that any activities undertaken by you will result in any specific sales, revenue growth, or other results.

All client examples described are presented as illustrations of how those clients have used IBM products and the results they may have achieved. Actual environmental costs and performance characteristics may vary by client.

Agenda

- Market trends
- Market evolution
- Feature highlights
- Key components
- Entry Points
- Cross Sell with IBM
- Sales Process
- Integrations
- Competition
- Packaging and pricing

Market trends

Today, enterprises are faced with multiple data challenges

More data

Exploding data growth

The aggregate volume of data stored is set to grow over **250%** in the next five years

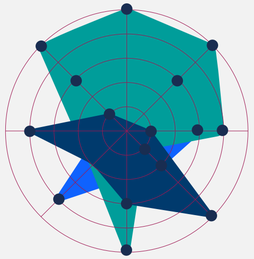

In multiple locations

Multiple locations, clouds, applications and silos

82% of enterprises are inhibited by data silos

In complex forms

Documents, images, video, and more

80% of time spent on data cleaning, integration and preparation

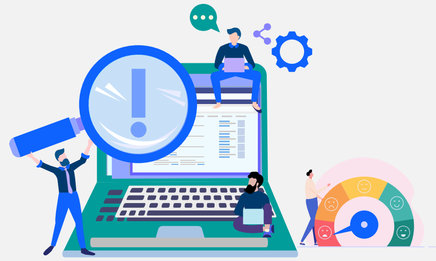

Poor quality

Stale and inconsistent

82% of enterprises say data quality is a barrier on their data integration projects

Data management market drivers

- Public cloud SaaS is overtaking on-prem.
- Transactional/operational systems represent the largest segment of the databases market, with $41B, 12% CAGR (‘21-’25).
- The warehouse & data lake market is $23B, 17% CAGR (‘21-’25) and almost all the growth is coming from public cloud SaaS and the data lake segment.

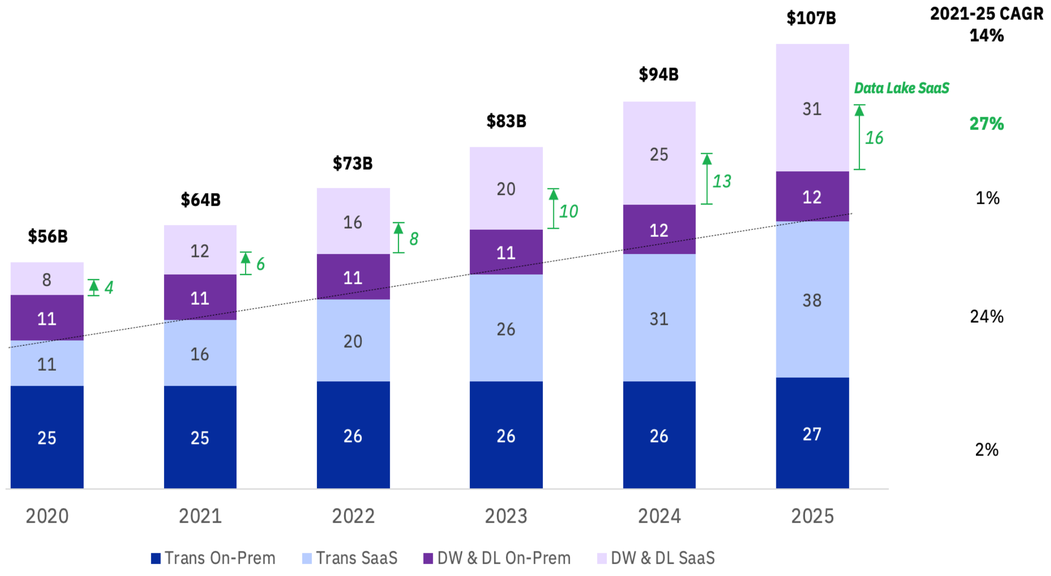

- Enterprises look to open source to:
  1. reduce license costs,
  2. leverage community innovation and support, and
  3. avoid vendor lock-in. Open source provides fit-for purpose technologies for diverse data types.

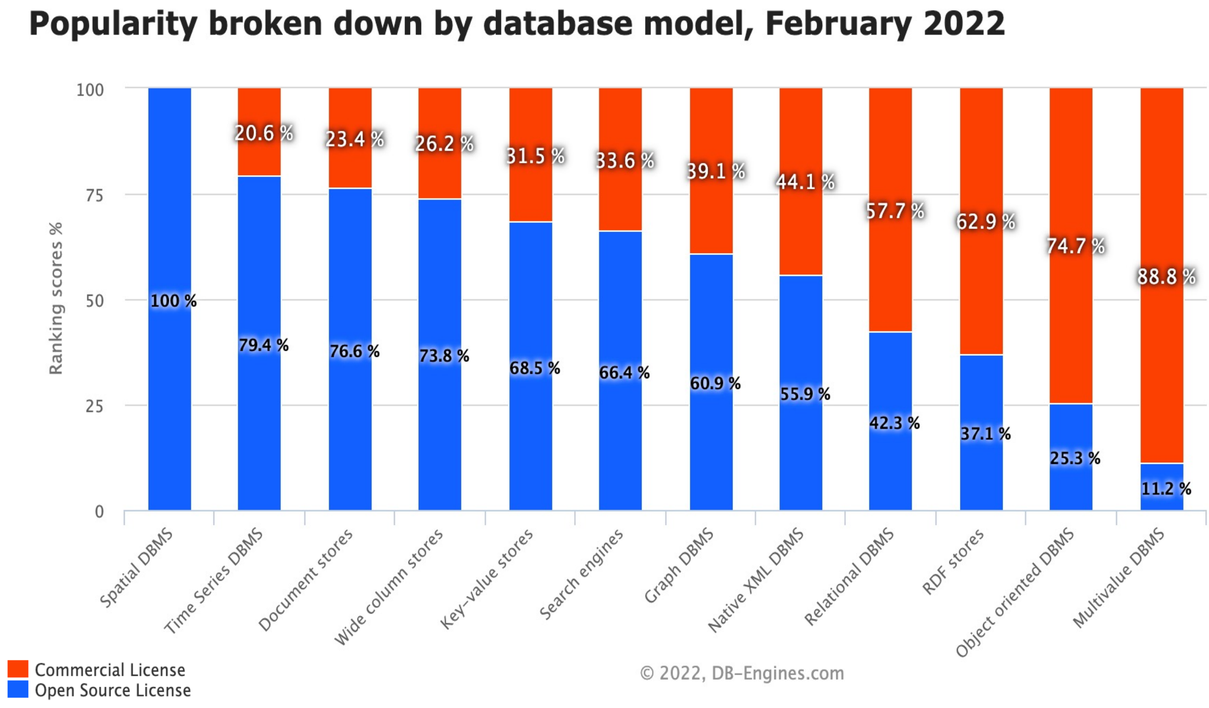

Market evolution

The data warehouse remains the center of analytics at most organizations

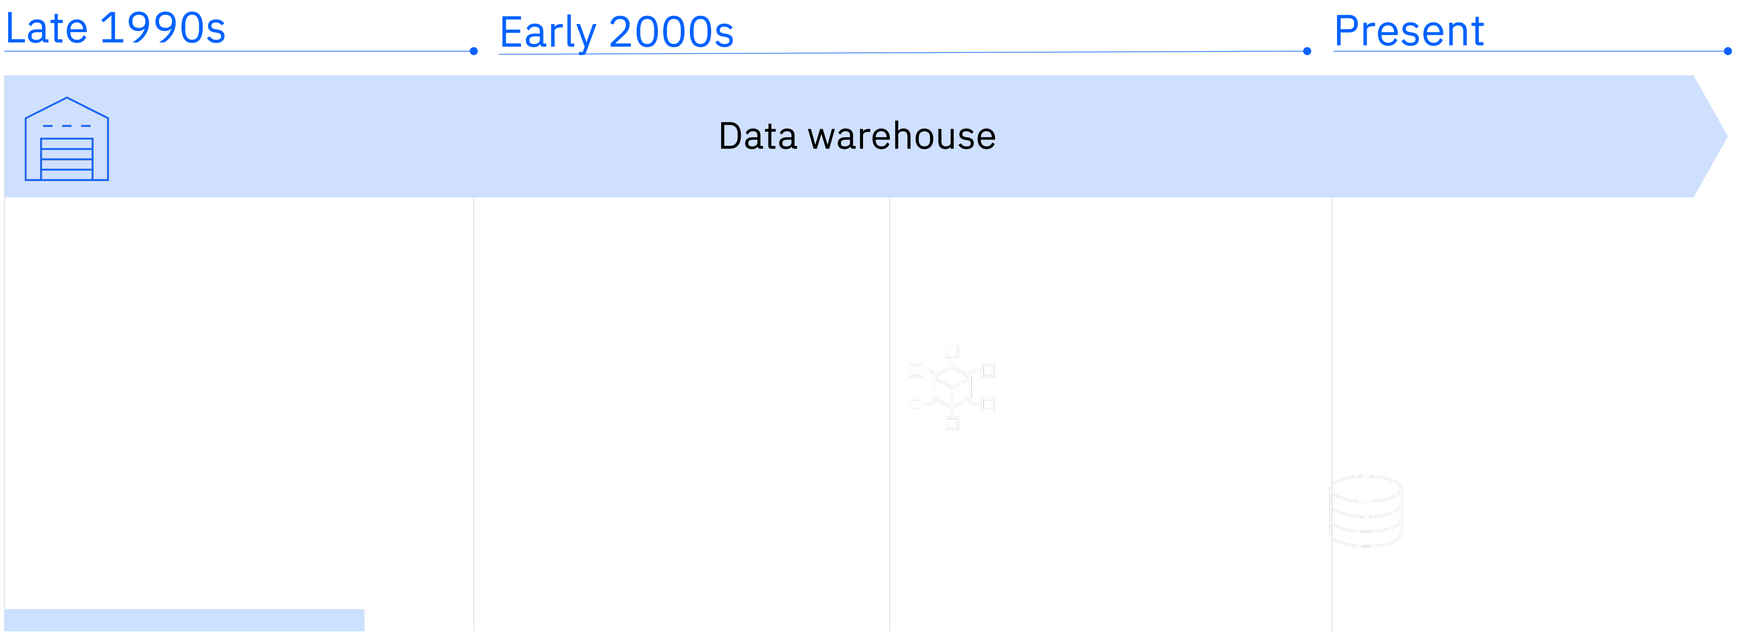

- Data warehouses emerged as the dominant method to analyze data
- Normalized and trusted data made it easy to analyze, however it’s an expensive choice
- Warehouses technology has evolved continuously to improve… from appliance form factors to in-memory technologies

The emergence of the data lake

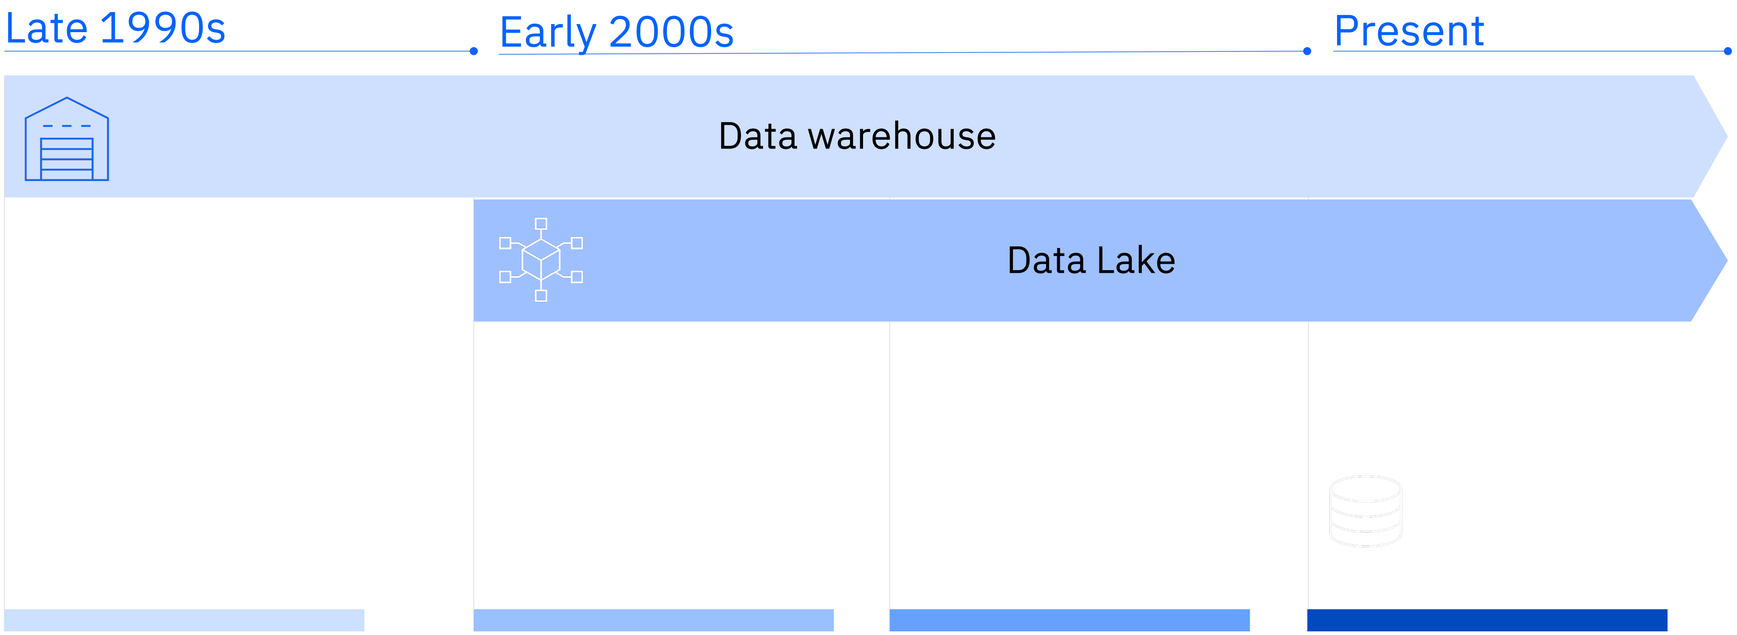

High up-front costs
Structured data only
ETL required
Vendor lock-in
Limited scalability

High complexity
Poor data quality
Limited performance
Expensive to maintain

- As volume, velocity, and variety of data grew, data lakes emerged as the new technology to replace data warehouses
- Data stored in raw and unstructured format = lower cost for large volumes of data
- Highly flexible and scalable
- Difficult to use & complex to maintain, and required a data scientist
- Ultimately, most data lakes failed and required a two-tier architecture

Cloud data warehouses evolved to address specific challenges of data warehouses

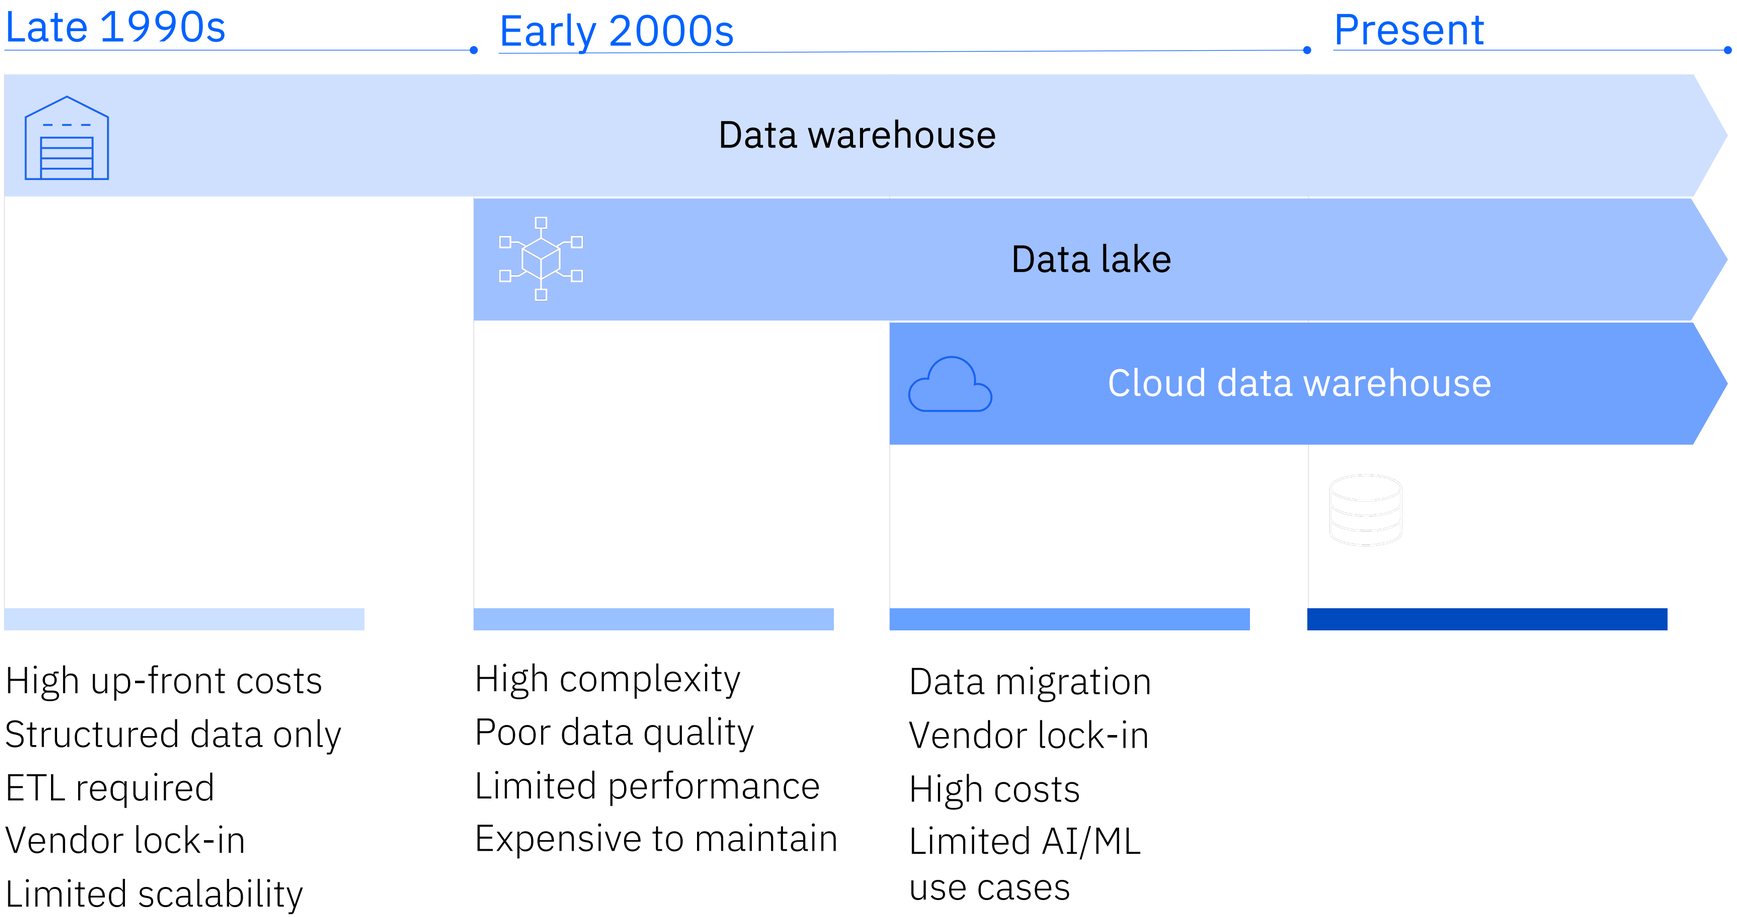

- Specifically, cloud data warehouses introduced the **separation of compute and storage**
- Addressed the scalability challenge with traditional warehouses – no data redistribution
- Ability to add more compute resources to the same data to solve the problem
- Easier to manage, however, much more expensive vs. on-prem warehouses

Data lakehouses are emerging technologies that solve for a new age of analytics

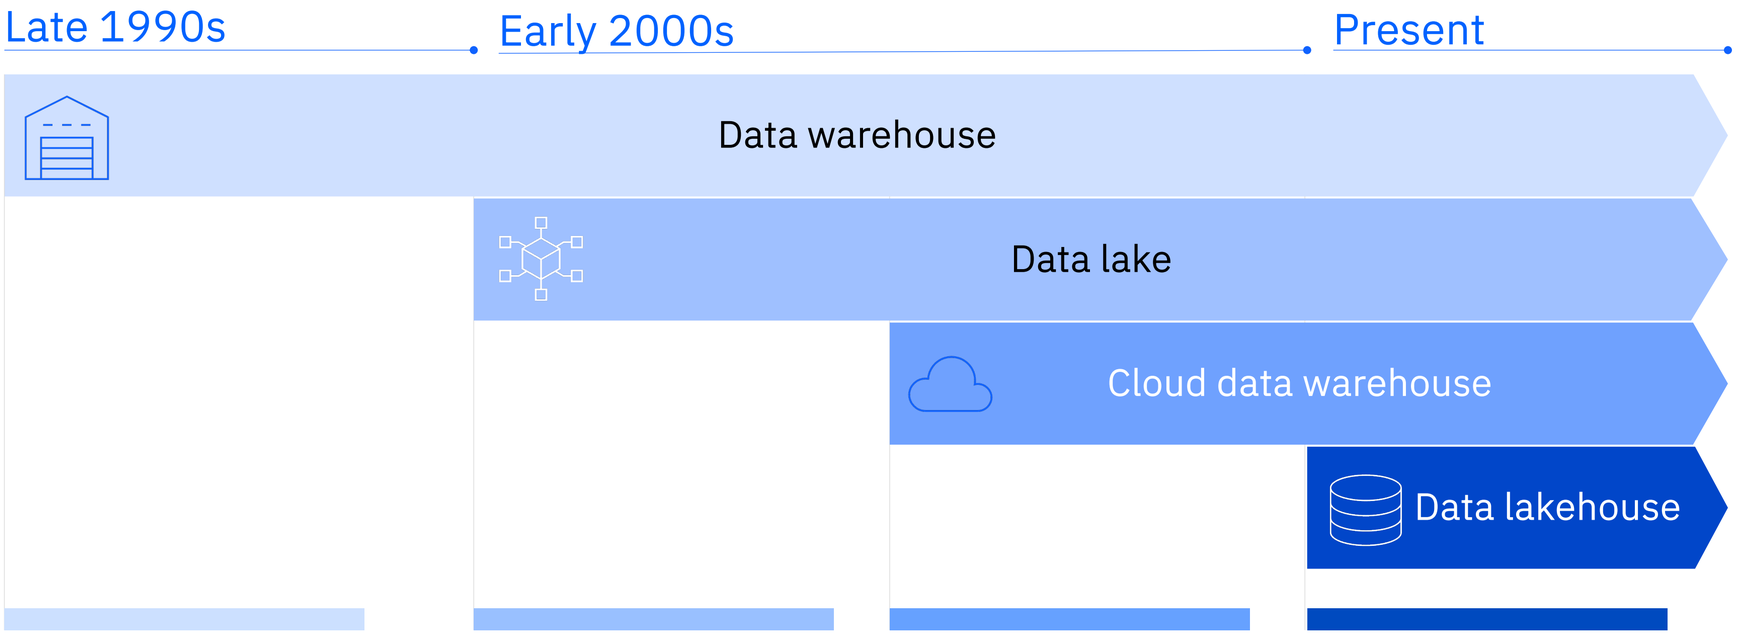

High up-front costs
Structured data only
ETL required
Vendor lock-in
Limited scalability

High complexity
Poor data quality
Limited performance
Expensive to maintain

Data migration
Vendor lock-in
High costs
Limited AI/ML use cases

Still early on tech curve
BI and AI in one place
Lower cost than warehouse
Open and flexible

- Most enterprises today require two-tier architectures – both a data lake and multiple warehouses
- Data is moved and replicated from lake to warehouse, and the warehouse is still the access layer for key data
- Data lakehouses combine the best of warehouses and data lakes
- Data warehouse engine plus data lake storage

Lakehouses are meant to be a new class of data store that combines the best of data warehouses and data lakes

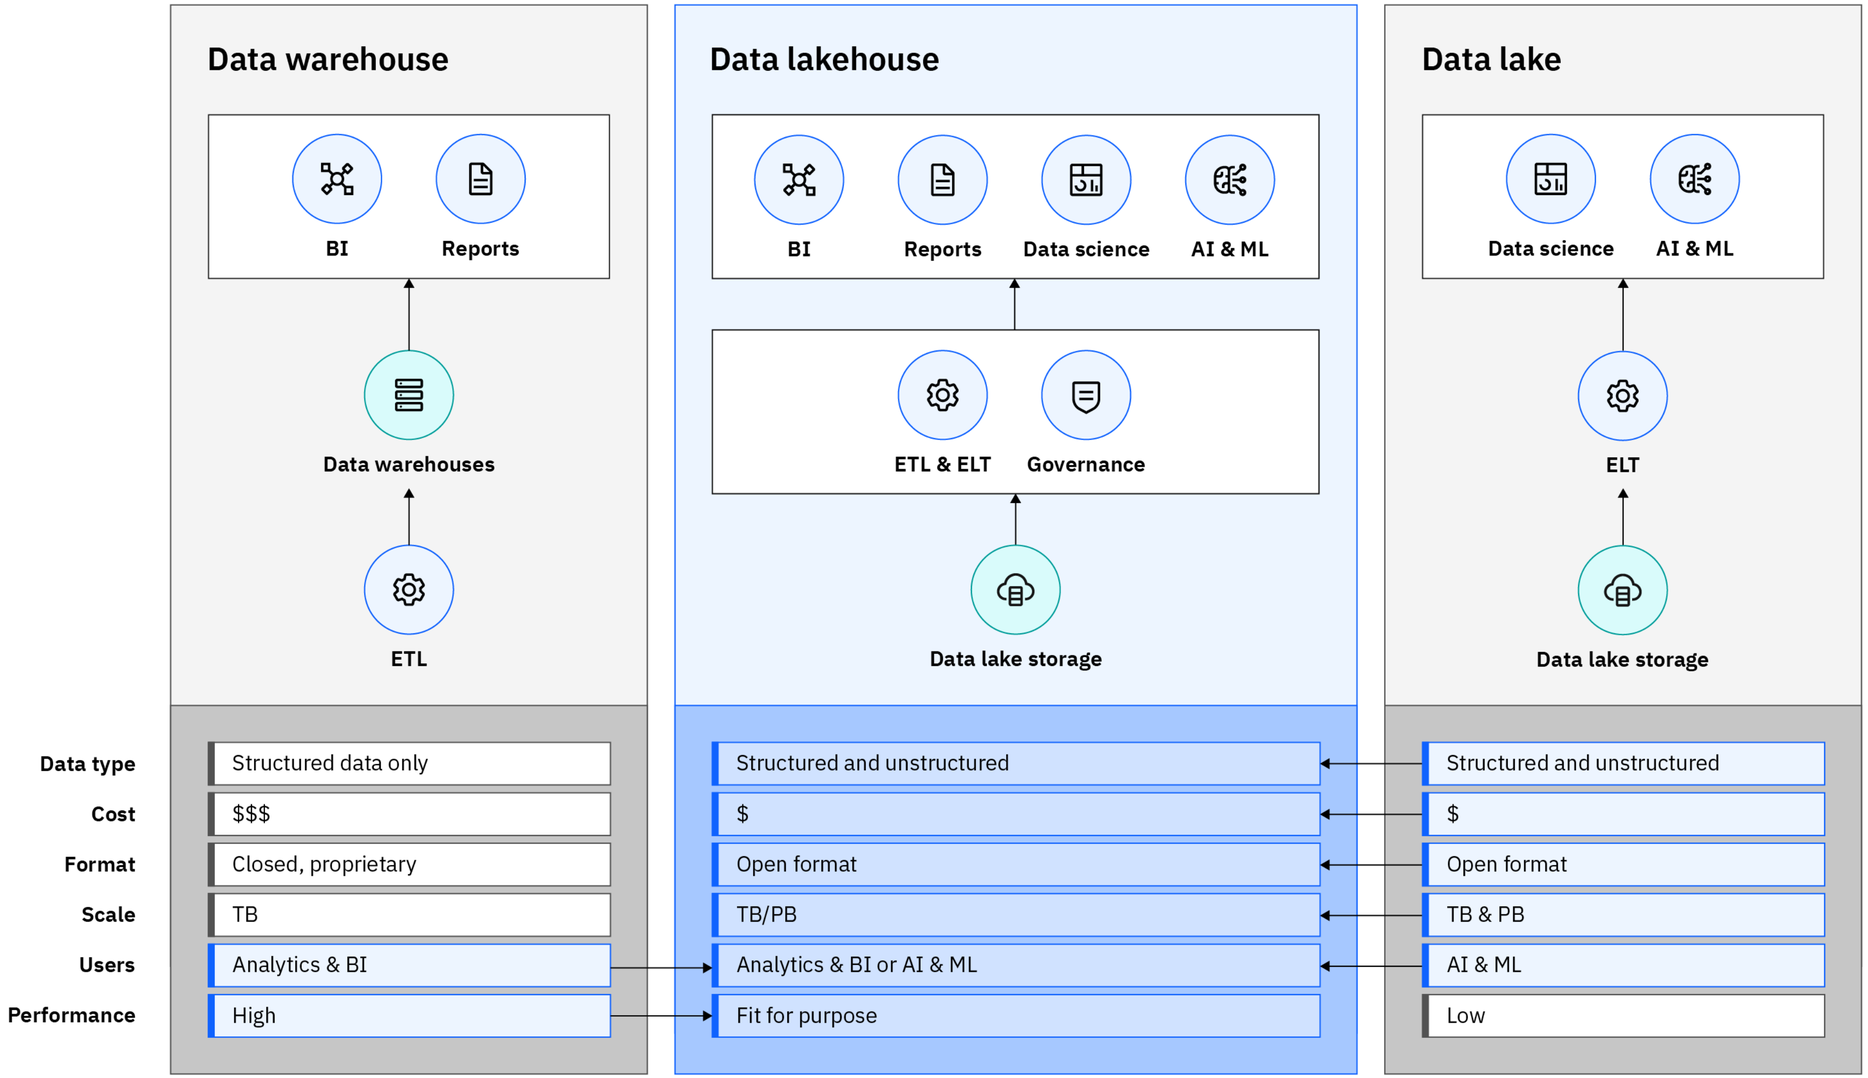

First generation lakehouses are still limited by their ability to address cost and complexity challenges:

- Single query engines set up to support limited workloads ... typically just BI or ML
- Typically deployed on cloud only with no support for multi-/hybrid-cloud deployments
- Minimal governance and metadata capabilities to deploy across the entire ecosystem

IBM watsonx.data feature highlights

IBM watsonx.data is the next evolution of current first-generation lakehouses

	First-generation lakehouse	Next generation lakehouse – IBM watsonx.data
	dremio, databricks, Starburst	presto, Apache Spark, ICEBERG, Governance and Metadata
Deployment:	Cloud or software deployment only	Hybrid cloud and multicloud deployment
Query engines:	Single engine with limited use case	Multiple fit-for-purpose engines
Governance:	Separate isolated functionality	Integrated and enterprise ready

IBM watsonx.data is the **only** lakehouse with **multiple query engines** allowing clients to **optimize costs and performance by pairing the right workload with the right engine**

Run all workloads from **a single pane of glass**, eliminating trade-offs with convenience while still improving cost and performance

Deploy anywhere **with full support for hybrid cloud and multicloud** environments

Shared metadata across multiple engines eliminates the need to re-catalog, accelerating time to value while ensuring governance and eliminating costly implementation efforts

IBM watsonx.data is designed to overcome the costs and complexities enterprises still face

The only open data store optimized for governed data and AI workloads across hybrid-cloud

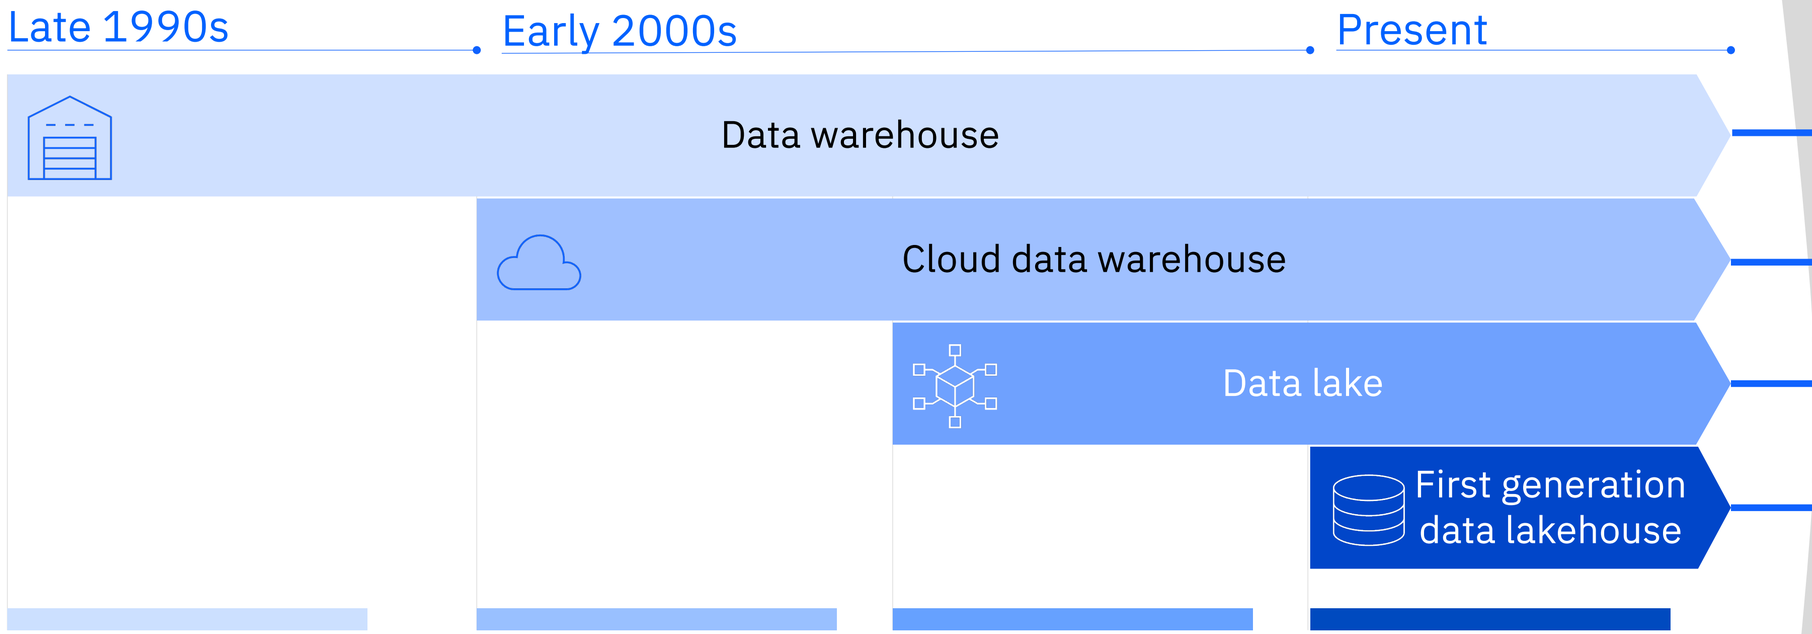

High up-front costs
Structured data only
ETL required
Vendor lock-in
Limited scalability

Data migration
Vendor lock-in
High costs
Limited AI/ML use cases

High complexity
Poor data quality
Limited performance
Expensive to maintain

Limited use cases
(BI only or AI/ML only)

Drive costs down by 50%, access 100% of the data, and get started in <10 minutes

Key components

Overview of the key components of IBM watsonx.data: multiple query engines, open table formats and built-in enterprise governance

A client's existing ecosystem

Databases

Data lakes

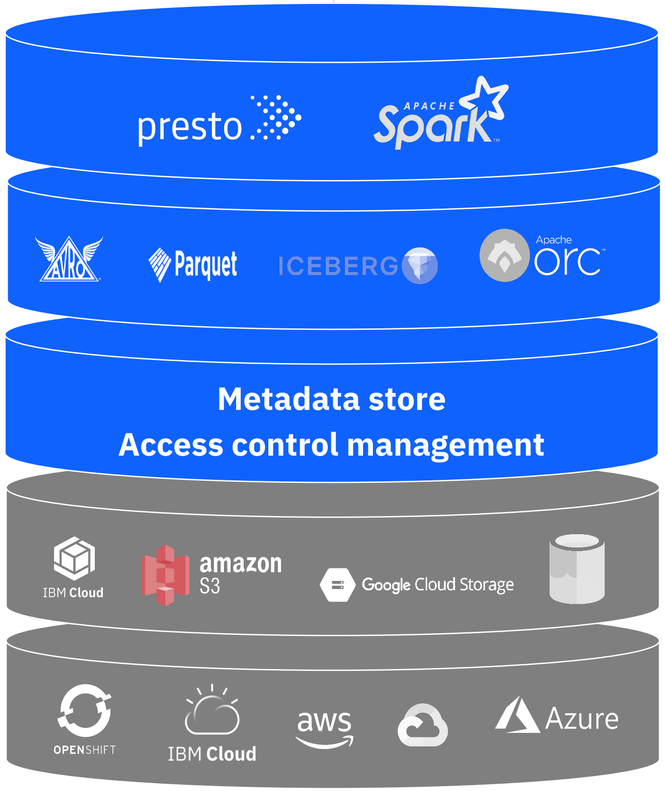

IBM watsonx.data

Core lakehouse functionality

Ecosystem infrastructure

Multiple engines such as Presto and Spark that provide **fast, reliable, and efficient processing of big data** at scale

Vendor agnostic open formats for analytic data sets, allowing different engines to access and share the same data, at the same time

Built-in governance that is compatible with existing solutions, including IBM WKC

Cost effective, simple object storage available across hybrid-cloud and multicloud environments

Hybrid-cloud deployments and workload portability across hyperscalers & on-prem with Red Hat OpenShift

- **Optimize workload costs and performance using multi-engine functionality**
- **Access 100% of the data across databases and data lakes**
- **Ensure governance and reduce time to insight with centralized metadata and access management**
- **Reduce storage costs and facilitate data ingest**
- **Deploy on any infrastructure and optimize available resources**

Three market entry points + use cases

1. Warehouse & Cost optimization narrative

- Competitive and IBM (Netezza + Db2)
- **Value prop:** Cost optimization and openness through the shared meta layer and fit-for-purpose engines
- Use case – Snowflake write-intensive workloads moving to Spark and/or Presto, thus reducing cost of Snowflake virtual data warehouses
- Use case – Db2 client wants to support both lake and warehouse workloads into a consolidated and modernized platform. Leverage Db2WoC SaaS gen3 with watsonx.data
- Use case- Offloading historical data from a warehouse that requires less frequent query access to save $$ on the warehouse.

2. Modernizing data lake narrative

- Modernizing storage architecture to facilitate shared metadata and fit-for-purpose engines
- **Value prop:** Cost optimization and openness through the shared meta layer and fit-for-purpose engines
- Use case – Move from HDFS to open source iceberg within a consolidated shared metadata layer
- Use case- Augment Cloudera to drive performance benefits and satisfy new workload requirements

3. Open source value add

- Replacing open source users of presto/trino
- **Value prop:** Enterprise support and functionality for OSS.
- Use case – Client is using presto today, move to watsonx.data for added functionality and support on opens source

Optimize costly warehouse workloads using fit-for-purpose engines that scale up/down automatically with a client's needs

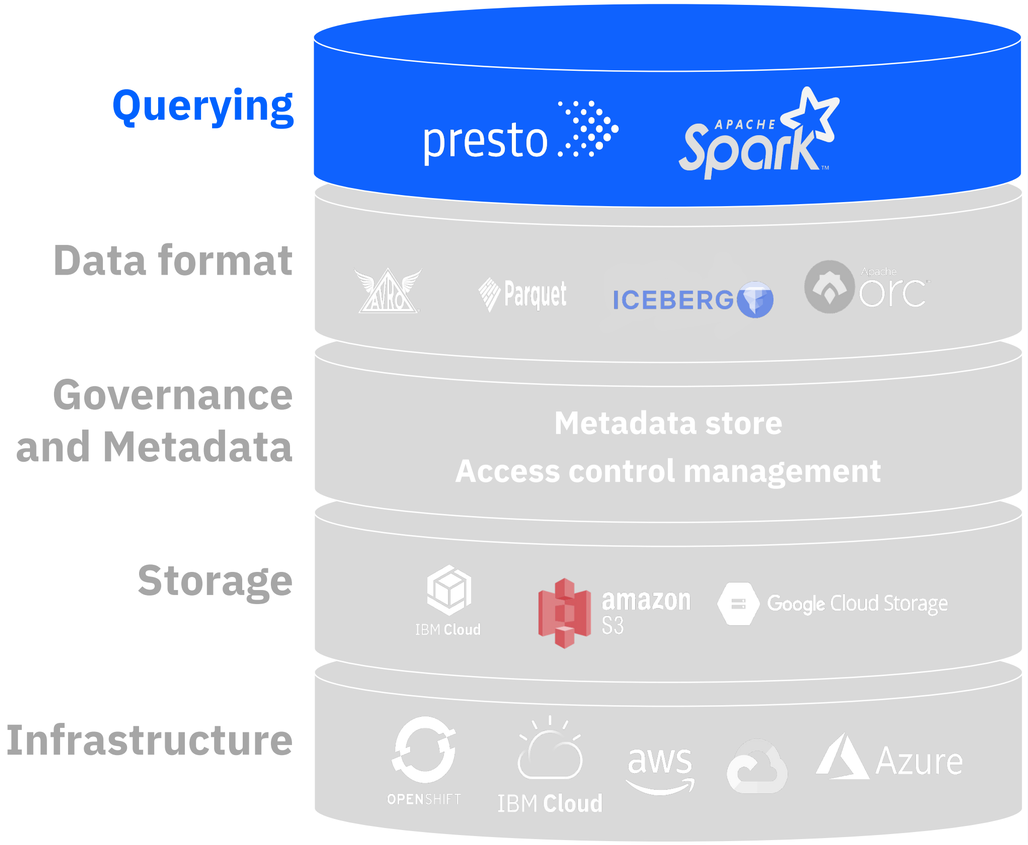

1 Utilize Iceberg data format and low-cost object storage to create a shared metadata layer that's accessible by multiple engines

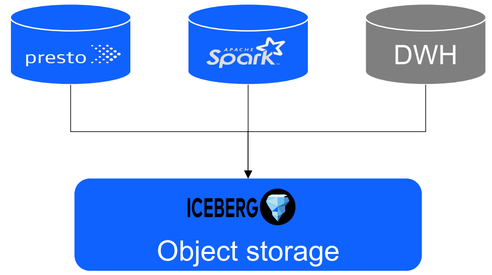

2 Run workloads on watsonx.data using fit-for-purpose compute and query engines

Use case	Query engine	Instance type
ELT/ETL	APACHE Spark	Compute
BI	presto	Cache
AI/ML	APACHE Spark	Compute

3 Take advantage of compute resources that scale up and down automatically based on workload requirements

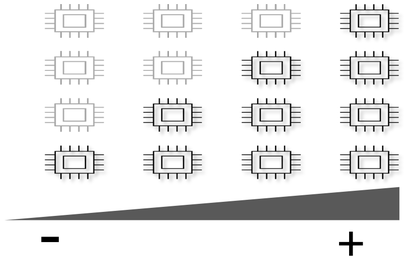

IBM watsonx.data reduces data warehousing costs by up to 50%*

Optimize workloads from data warehouses by taking advantage of low-cost object storage and fit-for-purpose query engines

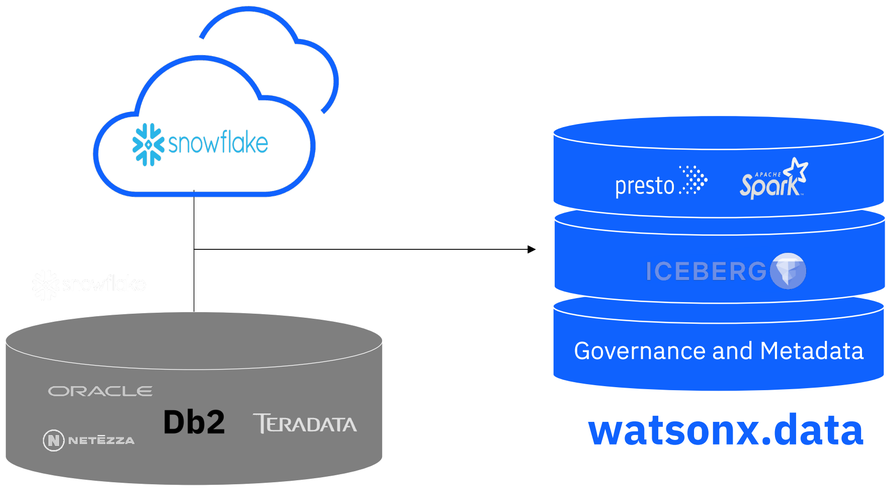

Utilize open table data formats and low-cost object storage

$$$$

Run workloads on the watsonx.data using fit-for-purpose query engines

Engine 1	Engine 2
-Reporting -BI	-Data science -ML

Improve infrastructure utilization with resources that scale up and down automatically based on workload requirements

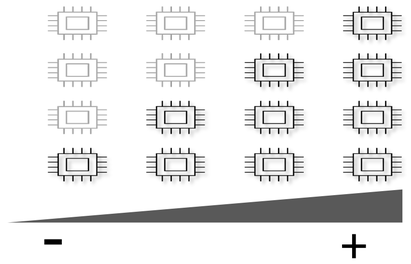

Up to 50%* savings compared to running workloads on cloud data warehouses

PrestoSQL to IBM watsonx.data

Enhance and operationalize open-source project such as PrestoSQL to enterprise ready offering like IBM watsonx.data

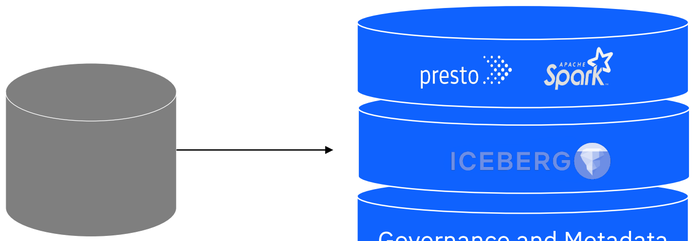

PrestoSQL / Trino

OSS Projects

watsonx.data

Only in Presto – *Not* in Trino

To understand more about Presto vs Trino, hear about the PrestoDB roadmap, watch this video from the Chair of the Technical Steering Committee and this one from Biswapesh Chattopadhyay, Tech Leader of Compute Infra at Facebook. Come join us in developing Presto, we welcome community participation. Below are some of the recent and current innovations.

Project Aria	PrestoDB can now push down entire expressions to the data source for some file formats like ORC. Blog Design
Project Presto Unlimited	Introduced exchange materialization to create temporary in-memory bucketed tables to use significantly less memory. PR Blog
Additional User Defined Functions	Support for dynamic SQL functions is now available in experimental mode. Docs
Presto-on-Spark	Presto on Spark Runs Presto code as a library within Spark executor. Facebook runs Presto for batch. Meetup talk Design Docs
Disaggregated Coordinator (a.k.a. Fireball)	Scale out the coordinator horizontally and revamp the RPC stack. In Beta in Q1 2021. PrestoCon Day Session (Mar 2021)
RaptorX Project	10x faster Presto with multiple level caching. PrestoCon Day Session (Mar 2021)

PrestoDB runs at Facebook

Trino does *not* run at Facebook

prestoDB

Meta

> "At Facebook alone, over a thousand employees use Presto, running several million queries and processing petabytes of data per day. After creating Presto we open sourced it to see if other companies were having the same issues and wanted to collaborate. It turns out many other companies were interested and so under The Linux Foundation, we believe the project can engage others and grow the community for the benefit of all."
>
> ***Kathy Kam, Head of Open Source at Facebook.***

Sample Discovery Questions

- What type of database are they using today
- Is cost an issue or starting to become an issue
- What are the use cases and how do they serve the business
- What are the required SLAs? Who are the users accessing? How fast do they need answers?
- Are there ETL workloads within the EDW, what are the time windows
- Is there historical data in the EDW that is not frequently accessed but still required to be there?
- Overall data size and instance size

IBM Cross Sell

What do I sell if my customer has:

- IBM analytics appliance?
- IBM Db2 Warehouse or Netezza?
- Db2 for z/OS?

IBM Modernize + Cross Sell

Customer Goal: Customer has an IBM on-premises solution and want to move to next generation of self-managed and SaaS offerings

Customer motivation:

- Looking for next generation features
- Limited compute or data storage resulting in offloading or expanding data infrastructure
- Difficulty scaling up and down for workloads
- New use cases: AI/ML, Real-time analytics, data engineering, data sharing

Challenges to Address:

- System migration compatibility
- Sizing management
- Cost controls on the cloud
- Access to all data across hybrid cloud without duplication/movement

Key Messages:

1. Modernize your database appliance with like for like compatibility and support new use cases
2. Optimize workloads for AI with fit for purpose engines and reduce warehouse costs by 50%, access all governed data across hybrid cloud

---

Warehousing with Db2:

- IBM® Integrated Analytics System (IIAS): EOS 2026
- Db2 Warehouse Software

Modernize appliance to Db2 Warehouse SaaS, Software or Cloud Rack →

Modernize software to Db2 Warehouse SaaS →

Db2 Warehouse SaaS
Fully managed service on AWS and IBM Cloud

Db2 Warehouse Software
(containerized software on cloud provider managed Kubernetes or OpenShift infrastructure)

Db2 Warehouse Cloud Rack
(containerized software on premises, OpenShift deployable on Power 10 or X86 infrastructures)

+

Warehousing with Netezza:

- Netezza Mako appliance
- Netezza CP4D Sys/Hammerhead

Modernize Mako to Netezza CP4D Sys (Hammerhead) or SaaS →

Modernize Hammerhead to Netezza SaaS →

Netezza CP4DSys
(containerized software provider managed Kubernetes or OpenShift infrastructure)

NPSaaS
Fully managed service on Azure and AWS

+

watsonx™

watsonx.ai
watsonx.data
watsonx.governance

IBM Db2 for z/OS Cross-sell

Customer Goal: Customer has IBM Db2 for z/OS and wants to access transactional data from mainframe for AI use cases

Customer motivation:

- Brand new use cases with mainframe transactional data: AI/ML, Real-time analytics, data engineering, data sharing

Challenges to Address:

- System migration compatibility

Key Messages z/OS and Db2 Warehouse:

- An integrated, optimized synchronization feature maintains currency between source Db2 for z/OS data on IBM Z and Db2 Warehouse targets
- Db2 Warehouse is a relational databases that delivers advanced data management and analytics capabilities

Key Messages z/OS and .data:

- Access and analyze data from mainframe systems in near-real time
- Consume less processor capacity with built-in synchronization
- Use the most up-to-date data from mainframe for machine learning models

---

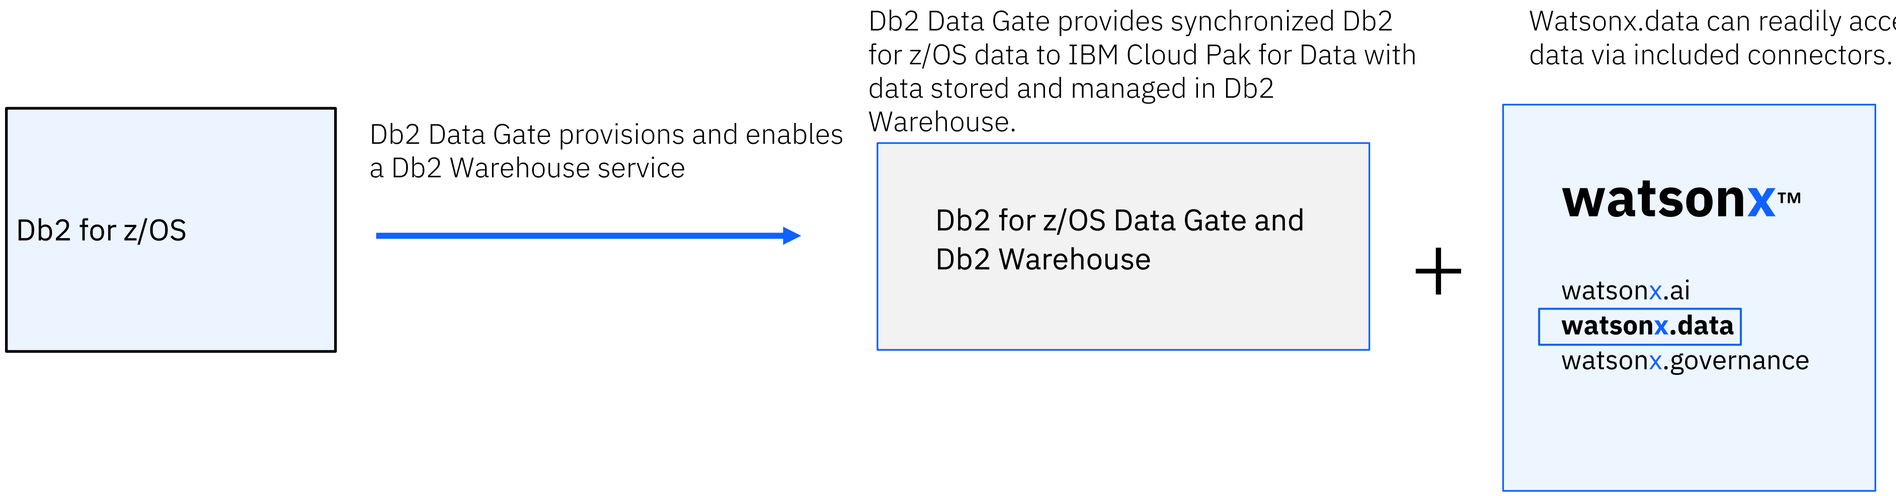

Sales Process View (non consulting led)

1. Discovery call/Demo
2. Workload discovery
3. Trial/POC- Client Engineering
4. Sizing/Pricing- Leverage input/Sales configurator
5. Sale- SQO and discount approval

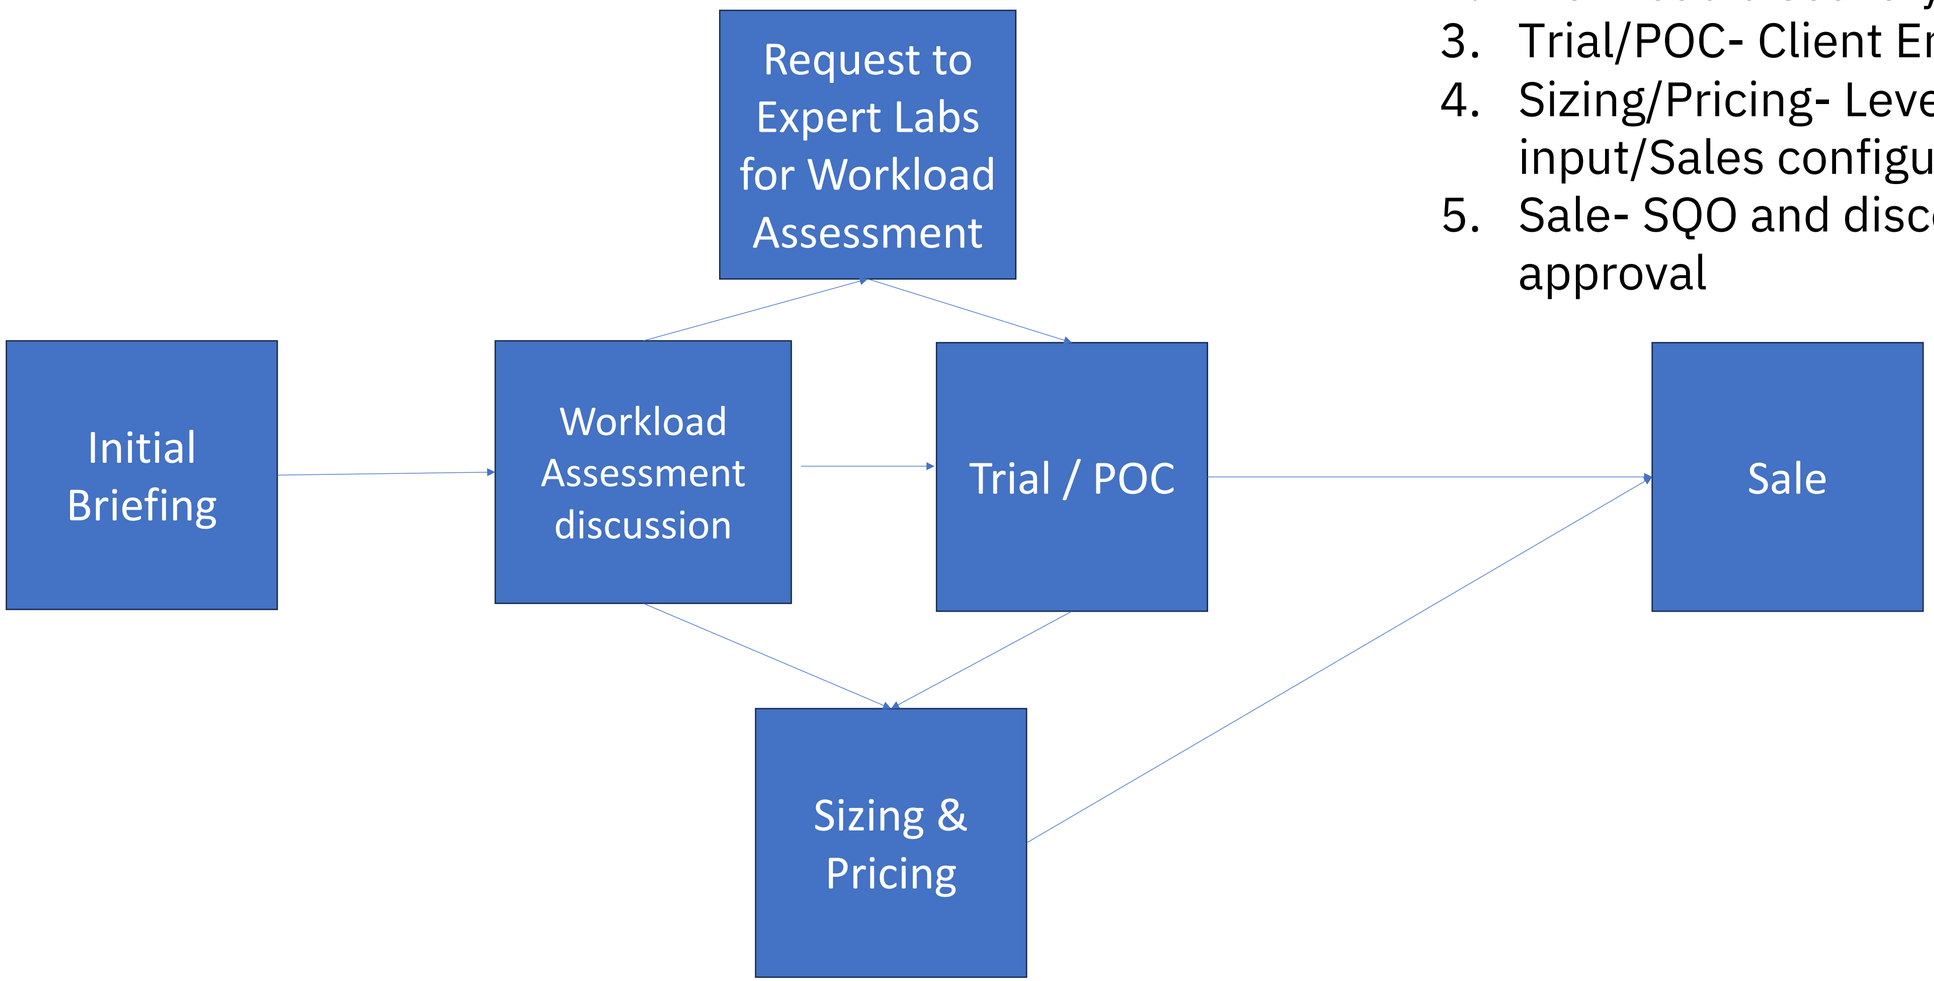

Integrations and IBM ecosystem for watsonx.data

Effortlessly populate watsonx.data with trusted data

Leverage best in class Data Ingestion and Observability

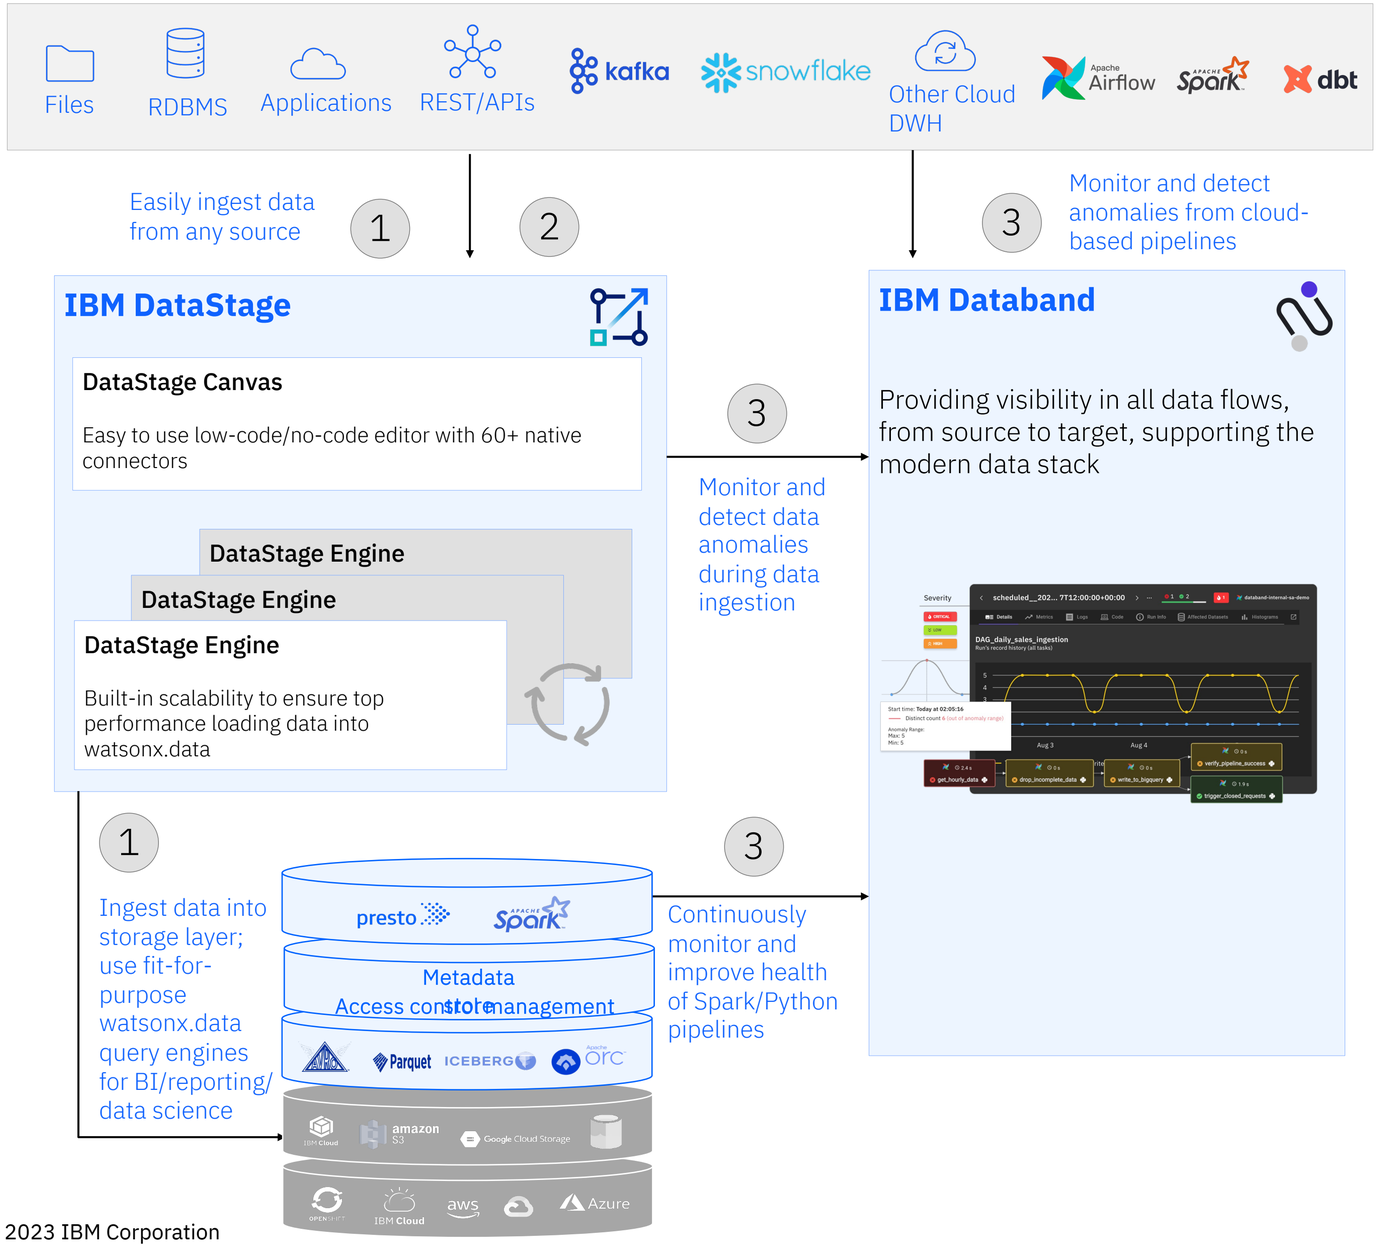

Data Pipelines with IBM DataStage

Easily build EL(T) pipelines with an intuitive visual design

1

Ingest data from any source

Leverage over 60+ native connectors to ingest data into watsonx.data from any type of source, ensuring top performance with built-in engine scalability

2

Reduce cost by offloading data from cloud data warehouses

Offload data from cloud data warehouses to enable shifting workloads like BI, reporting, or data science to fit-for-purpose query engines

Data Observability with IBM Databand

Continuously detect and resolve data quality incidents

3

Monitor, detect, and resolve data quality incidents

Monitor and improve the health of DataStage, Spark, or Python pipeline workloads running on watsonx.data

Detect data anomalies and accelerate issue resolution

The integrated IBM watsonx.data ecosystem for maximum workload coverage and optimal price-performance

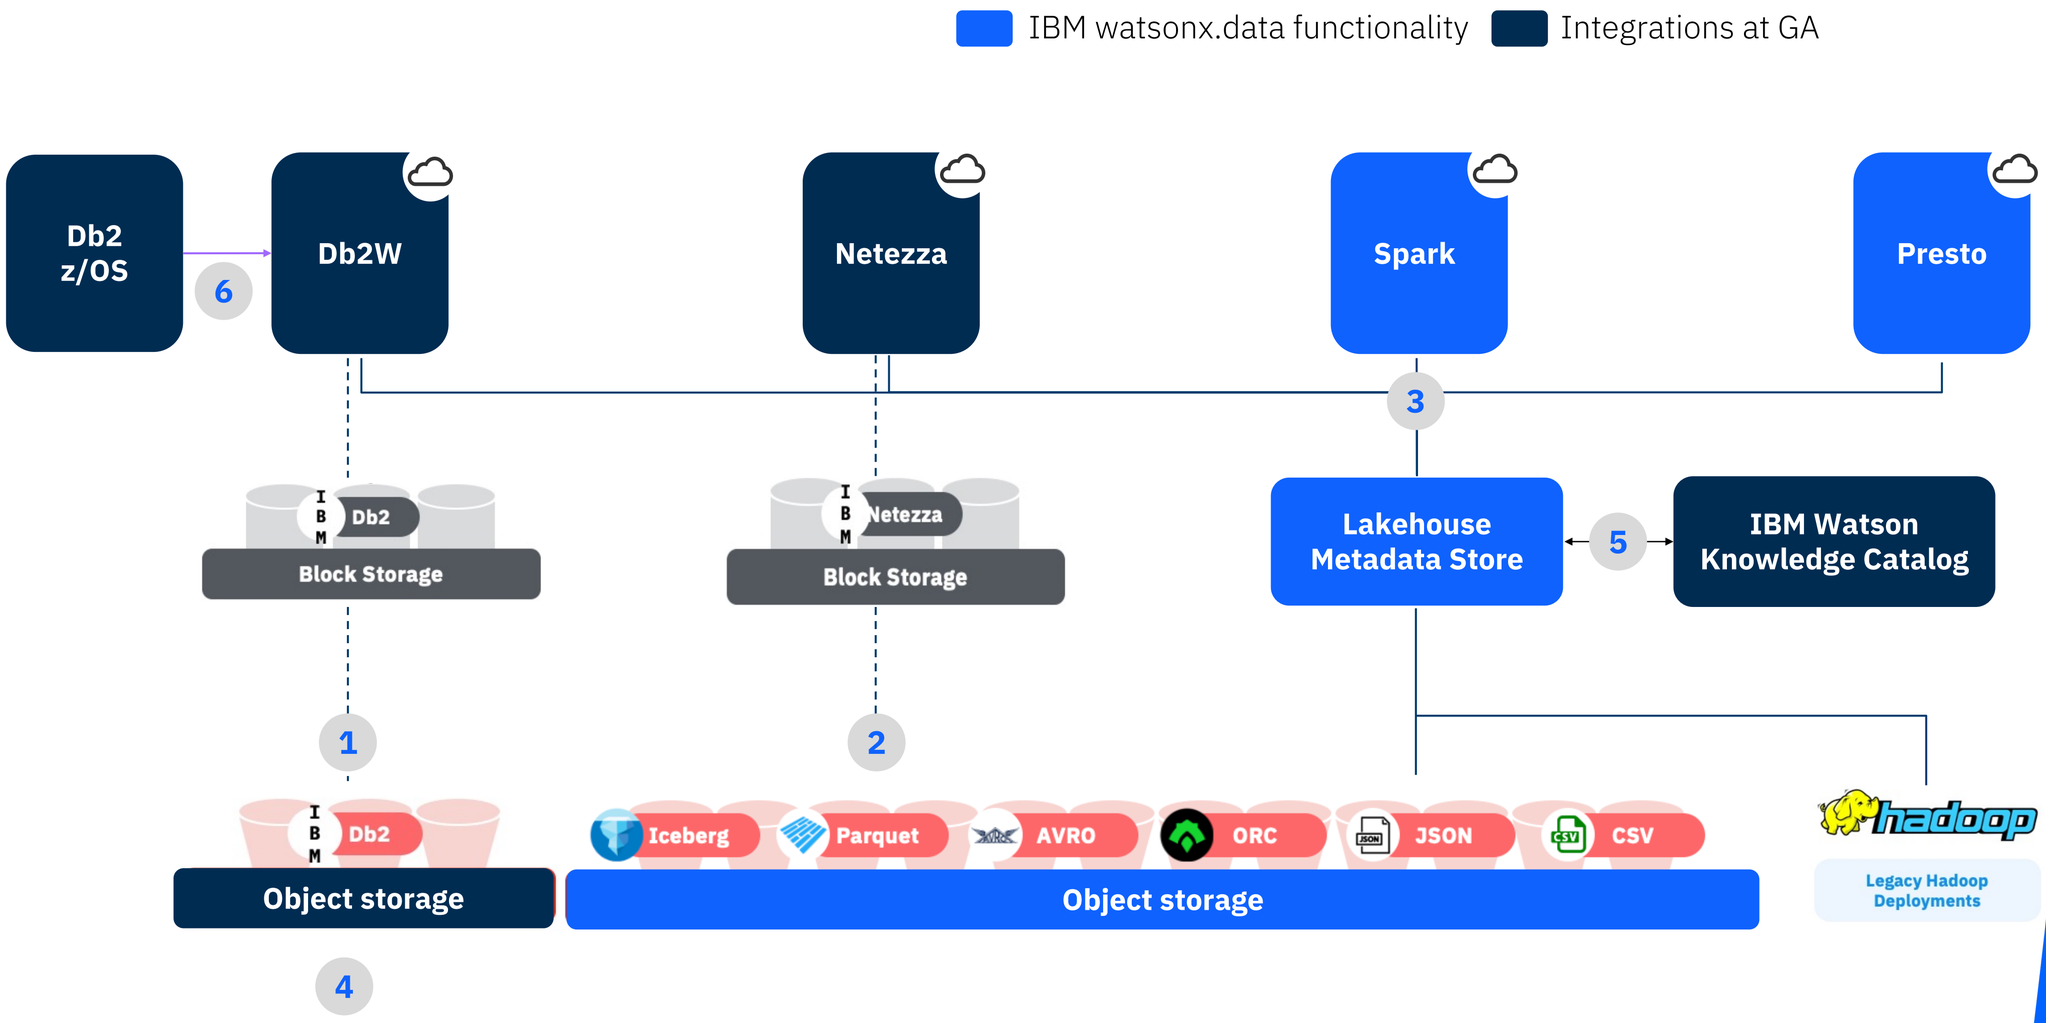

1. Warehouses can access data in the lakehouse
2. Easily Promote data between the warehouse and lakehouse
3. Query routing service, multiple engines can access same data lake data
4. The lakehouse can access data residing in Db2/Netezza
5. WKC policies enforced by the lakehouse via metadata service
6. Analyze Z data easily and securely with Db2 for z/OS Data Gate

IBM Databases + watsonx.data Cross Sell Motion Sales Kit (all on Seismic)

Relevant Client Facing Asset	Type	Description	Link
Watsonx.data 1 pager	1 pager	Client facing documents can be emailed to clients to increase customer product understanding/knowledge on watsonx.data	Live
Watsonx.data Solution Brief	Solution Brief		Live
Watsonx.data client deck	201 Client Deck	This deck includes an overview of watsonx.data, how IBM can provide value-add to transform a customers data management architecture.	Live
Watsonx.data 3 touch email template (w/existing customers)	Email	Emails for sellers to reach into accounts and identify watsonx.data customer opportunities.	Live
Db2 Warehouse + watsonx.data 1 pager	1 pager	Client facing document summarizing benefits of the next generation of Db2 Warehouse + watsonx.data	Live
Db2 Warehouse Solution Brief	Solution Brief	Client facing document detailing capabilities of the next generation of Db2 Warehouse and watsonx.data integration	Live
Netezza + watsonx.data 1 pager	1 pager	Client facing document summarizing benefits of Netezza aaS on AWS + watsonx.data	Live
Netezza Solution Brief	Solution Brief	Client facing document detailing capabilities of NPSaaS and watsonx.data integration	Live
PDOA/IIAS appliance + Db2 Warehouse software modernization 1 Pagers	1 pager	Client facing document summarizing benefits	PDOA/IIAS: Live Db2 Warehouse Software: Live
Watsonx.data Demos	Demo	Demo videos: Getting started, fit for purpose query engines, single copy of data, powered by generative AI	Live on Seismic
Db2 Warehouse Next Generation SaaS Demo	Demo	Demo video showcasing next generation of Db2 Warehouse with watsonx.data integration	Live

Competitors and objection handling

Competitive landscape

Data lakehouse competitors

Details	databricks	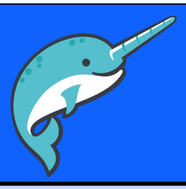	Starburst	
Deployment options	Public cloud only	Public cloud & on-premises	Public cloud & on-premises	AWS only
Query engines	• Apache Spark • Photon	• Dremio Sonar (proprietary)	• Starburst (Trino-based)	• Apache Spark • Amazon (Trino-based)
Open table format support	• Delta Lake	• Apache Iceberg	• Apache Iceberg • Delta Lake	• Apache Iceberg (Parquet only)

Other competitors

Details	snowflake	
Deployment options	Public cloud only	AWS only
Query engines	• Snowflake	• Amazon Redshift
Open table format support	• Supports Apache Iceberg but primarily uses Cloud Object Storage (COS)	• None, uses Amazon S3 files defined as external tables

Packaging and pricing

watsonx Pricing (for 7/7 GA)

watsonx.data (SaaS)

Consumption-based charges, driven off 3 Price metrics*

- Quantity of nodes deployed
- Duration of nodes deployed
- Support services deployed (e.g., Access Control, Metastore)

*watsonx.data only available for Standard and Premium Tier Customers

Pricing

	Price (USD) / Hour
Cache coordinator node per hour	$2.80
Cache worker node cost per hour	$2.80
Compute coordinator node per hour	$6.50
Compute worker node cost per hour	$6.50
Supporting services per hour	$3.00

watsonx.data (On-Prem)

- Price Metric: VPCs (16 VPCs/node) @ $7500/VPC
- Minimum of 10 nodes (160 VPCs) recommended
- Production versions available in Perpetual, Subscription and Monthly Licenses. Options for Non-production, reserved, and various support tiers also available.

Pricing

License Type	Price (USD) / VPC
Perpetual	$7,500
Subscription	$250 / month
Monthly	$312.50 /month

On-prem Parts are listed here

Leverage the sales configurator for sizing- SW and SaaS

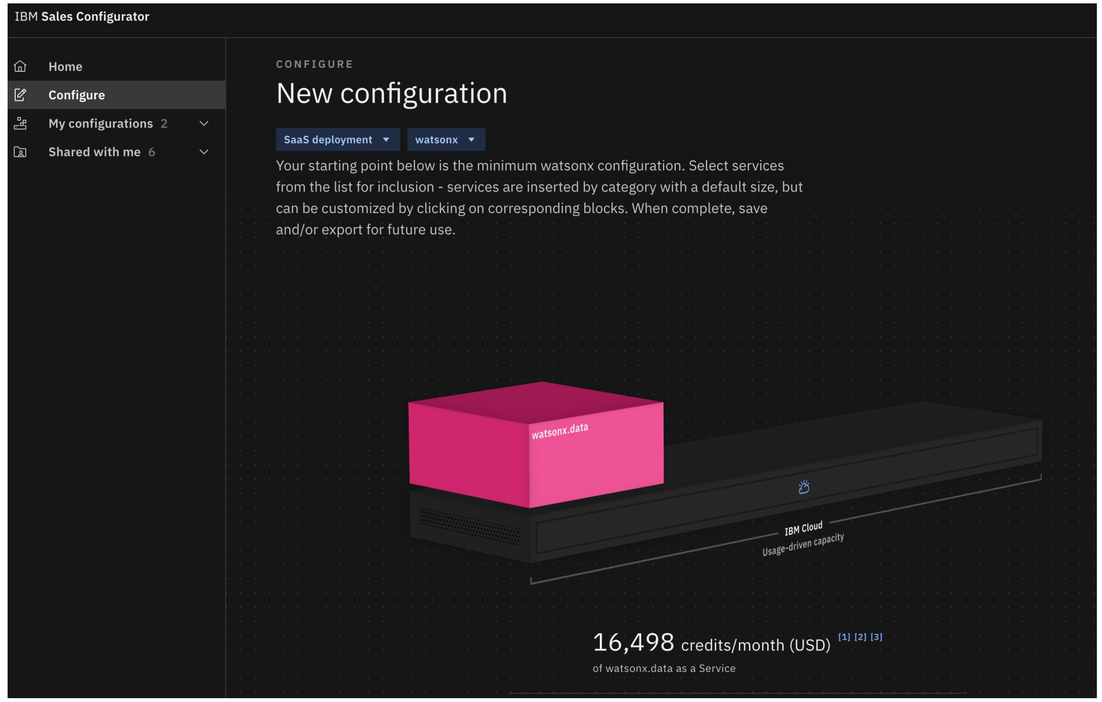

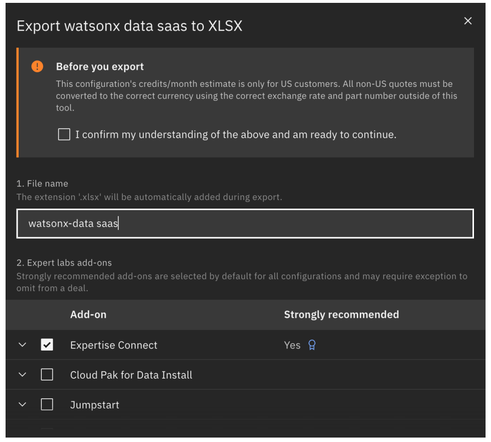

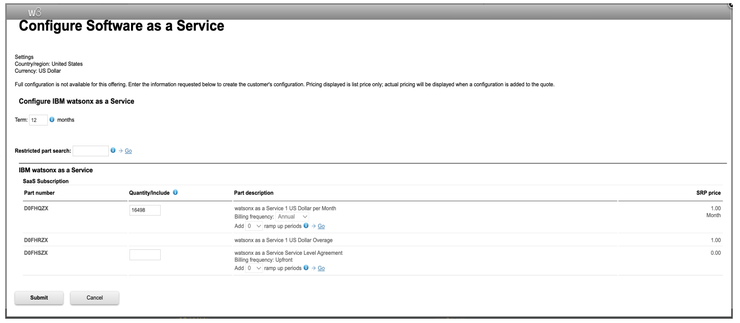

Note: Recommend engaging Expert Labs Services and Client Engineering to accelerate productive use

watsonx.data SaaS Pricing – What you need to know

Consumption-based charges, driven off 3 Price metrics*

- Quantity of nodes deployed
- Duration of nodes deployed
- Support services deployed (e.g., Access Control, Metastore)

*watsonx.data only available for Standard and Premium Tier Customers

Consumption measured in Resource Units (RU). 1 RU = $1 USD

Pricing	**RU / Hour**
Cache node cost per hour	$2.80
Compute node cost per hour	$6.50
Supporting services per hour	$3.00

Key Takeaways

- Usage can occur on IBM Cloud or AWS infrastructure
- Pricing depends on type of nodes, number of nodes used, and usage of nodes
  - AWS infrastructure provides two node types (Cache and Compute Optimized), IBM Cloud infrastructure has one node type (Cached Optimized). Prices are the same for AWS and IBM Cloud
  - Number of nodes relates to T-shirt sizes.
  - Usage relates to uptime of services.
- Pricing is measured in Resource Units. 1 RU = $1 USD

Estimated T-Shirt Sizes for IBM Cloud

Sizing	Nodes	Total vCPU	Total RAM (GiB)	Resource Unit per Month*
Starter	2 nodes	32 vCPU	256 GB	•Using service 100% of the time: 6278 •Using service 70% of the time: 5052 •Using service 35% of the time: 3621
Small	4 nodes	64 vCPU	512 GB	•Using service 100% of the time: 10366 •Using service 70% of the time: 7913 •Using service 35% of the time: 5052
Medium	7 nodes	112 vCPU	896 GB	•Using service 100% of the time: 16498 •Using service 70% of the time: 12205 •Using service 35% of the time: 7198
Large	12 nodes	192 vCPU	1536 GB	•Using service 100% of the time: 28762 •Using service 70% of the time: 20790 •Using service 35% of the time: 11490

*Resource Unit Values in the table can be converted to US Dollar value at 1:1 Ratio

Standard (on-Prem) T-shirt Sizing

T-shirt size for a single cluster with performance characteristic: If customer wants multiple identical environments for different departments or Dev+Prod, multiple across
This assumes a worst-case scenario - 100% of your data is in active memory. **Note**: These sizes below are each for 1 cluster

	Description	License (VPC) for 1 cluster of this profile	Base Entitlement (UI,HMS, Etc)	Total Entitlement	Total HW requirements
Small	1 Coordinator node, 3 worker nodes Each node assumes 16vCPU, 128Gb of memory, and 2 x 1900GB NVMe, up to 10 GbE. Example of estimated workload supported by the configuration: • 1 to 5 queries processing 100GB → 400 GB expanded/uncompressed active data in memory or up to 50 queries processing up to 8GB each • 100-200GB Active data compressed on disk required to be processed.	64	12	76	76 vCPU Total Memory:608Gb
Medium	1 Coordinator node, 9 worker Each node assumes 16vCPU, 128Gb of memory, and 2 x 1900GB NVMe, up to 10 GbE. Example of estimated workload supported by the configuration: • 1 to 5 queries processing 250GB → 1125 GB expanded/uncompressed active data in memory or up to 50 queries processing up to 22GB each • 250-500GB Active data compressed on disk required to be processed.	160	24	184	184vCPU Memory: 1472Gb
Large	1 Coordinator node 19 worker Each node assumes 16vCPU, 128Gb of memory, and 2 x 1900GB NVMe, up to 10 GbE. Example of estimated workload supported by the configuration: • 1 to 5 queries processing 500GB → 2250 GB expanded/uncompressed active data in memory or up to 50 queries processing up to 45GB each • 500-1000GB Active data compressed on disk required to be processed.	320	48	368	368vCPU Memory: 2944Gb
X-Large	70 Nodes to configure as desired (i.e 10 clusters of 1 head node and 6 worker nodes or 1 large cluster with 70 nodes). Each node assumes 16vCPU, 128Gb of memory, and 2 x 1900GB NVMe, up to 10 GbE. Example of estimated workload supported by the configuration: • 1 to 5 queries processing 1750GB → 7875 GB expanded/uncompressed active data in memory or up to 175 queries processing up to 45GB each • 1750-3500GB Active data compressed on disk required to be processed.	1120	48	1168	1168vCPU Memory: 9344Gb

Technology Expert Labs **watsonx**.data Companion Services

Install Offering
(Non-SaaS only)

2 Weeks

- ✓ Services offering to deploy and configure a **watsonx**.data environment & integrate with the client IT environment

1. Install and configure OpenShift (non-SaaS only).
2. Install and configure **watsonx**.data service

- ❖ Outcome: Scalable **watsonx**.data environment ready for client workload

Build Offering

3 Weeks

- ✓ Services offering to implement one supported use case with **watsonx**.data. Suitable for on-prem or saas.

1. Integrate with client infrastructure
2. Implementation & validation
3. User Acceptance testing & go live support

- ❖ Outcome: Successful optimization of client workload with **watsonx**.data

Expertise Connect

1 Year

- ✓ Subscription-based service that provides access to deep technical expertise in design, deployment, operations, and transformation while building your client's capacity and skills

- ❖ Expertise Connect is the insurance policy that helps clients succeed with IBM technologies and allows them to ***move forward*** in their IBM journey.

Custom Services

- ✓ Tailored services to help a customer augment their existing Data Management systems with **watsonx**.data.

- Planning
- Solution Design
- Installation
- Workload Optimization
- User Acceptance Testing
- Go-Live Support
- Enablement

Engage Solution Engineering

Client path to **watsonx.data** adoption through Technology Expert Labs →

1

Attach pre-scoped offerings to all new license sales

2

Expand with Expertise Connect

3

Engage Expert Labs Services Sellers early in the Software Sale

Expert Labs Contacts

Jennifer Wales
Services Product Manager, Data and AI

Ted Trask
Program Director, WW Data & AI Services Sales Leader

Pradeep Kutty
Business Unit Executive – Solution Engineering

Suzanne Golledge WW Principal Delivery Practice Leader

Seller Locator Tool

Slack: #ask-expert-labs

Seismic Offering Page	Retail Price	Duration	Transact by SaaS	Transact by Non-SaaS
Build	$55K USD	3 W	Qty 3 x #D06U8ZX Valid July only*	Qty 3 x #D06U8ZX Valid July only*
Install (non-SaaS only)	$36K USD	2 W	--------	Qty 2 x #D06W3ZX Valid July only*
Expertise Connect	$90K USD	6M	Qty 1 x D0676ZX	

*New Watsonx parts coming late July – see Seismic Offering page for latest parts